



6CS030 Lecture 8

SQL on Hadoop
Hadoop Stack '**Zoo**':
 Apache Spark
 Apache HBase
 Apache Pig
 Apache Hive

SQL on Hadoop

- MapReduce is very complex when compared to SQL
- Need for a more database-like setup on top of Hadoop
- Various projects can be used on top of Hadoop
 - See <http://hadoop.apache.org/> for a list.
- Sometimes referred to as a “Zoo”
- **ZooKeeper** provides a high-performance coordination service for distributed applications

SQL on Hadoop

- HBase

- <https://hbase.apache.org>

- Pig

- <https://pig.apache.org/>

- Hive

- <https://hive.apache.org/>

- Spark

- <https://spark.apache.org/>



HBase



- First Hadoop database inspired by Google's Bigtable
- Runs on top of HDFS
- NoSQL like data storage platform
 - No typed columns, triggers, advanced query capabilities, etc.
- Offers a simplified structure and query language in a way that is highly scalable and can tackle large volumes

HBase



- Similar to RDBMS, HBase organizes data in tables with rows and columns
- HBase table consists of multiple rows
- A row consists of a row key and one or more columns with values associated with them
- Rows in a table are sorted alphabetically by the row key

HBase



- Each column in HBase is denoted by a column family and qualifier (separated by a colon, ':')
- A column family physically co-locates a set of columns and their values
- Every row has the same column families, but not all column families need to have a value per row
- Each cell in a table is hence defined by a combination of the row key, column family and column qualifier, and a timestamp

HBase



```
Creating a HBASE table with Column Family Schema  
  
hbase(main):001:0> create 'users', 'name', 'email'  
  
0 row(s) in 2.8350 seconds  
=> Hbase::Table - users
```

HBase users Table (Conceptual View)

Row Key	name:first	name:last	email:work	email:personal
user1	John	Doe	john.doe@company.com	johnny123@gmail.com
user2	Alice	Smith	alice@tech.com	alice.smith@yahoo.com
user3	Bob	Brown	bob@startup.io	bob_b@hotmail.com

HBase



```

Description Result

{
  "users": {
    "column_families": {
      "email": {
        "NAME": "email",
        "BLOOMFILTER": "ROW",
        "VERSIONS": "1",
        "IN_MEMORY": "false",
        "KEEP_DELETED_CELLS": "FALSE",
        "DATA_BLOCK_ENCODING": "NONE",
        "TTL": "FOREVER",
        "COMPRESSION": "NONE",
        "MIN_VERSIONS": "0",
        "BLOCKCACHE": "true",
        "BLOCKSIZE": "65536",
        "REPLICATION_SCOPE": "0"
      },
      "name": {
        "NAME": "name",
        "BLOOMFILTER": "ROW",
        "VERSIONS": "1",
        "IN_MEMORY": "false",
        "KEEP_DELETED_CELLS": "FALSE",
        "DATA_BLOCK_ENCODING": "NONE",
        "TTL": "FOREVER",
        "COMPRESSION": "NONE",
        "MIN_VERSIONS": "0",
        "BLOCKCACHE": "true",
        "BLOCKSIZE": "65536",
        "REPLICATION_SCOPE": "0"
      }
    }
  }
}

```

```

Describing The Created Table

hbase(main):002:0> describe 'users'

```

```

Description Result

Table users is ENABLED
users
COLUMN FAMILIES DESCRIPTION

```


HBase



- Now lets, build up on the initial contextual example : HBase table to store and query users
- The row key will be the user id
- column families:qualifiers
 - name:first
 - name:last
 - email (without a qualifier)

HBase



Upserting Value to A HBASE Table

```
put 'users', 'seppe', 'name:last', 'vanden Broucke'  
put 'users', 'seppe', 'email:work', 'seppe.vandenbroucke@kuleuven'
```

Scanning A HBASE Table

ROW	COLUMN+CELL		
seppe	column=email:work,	timestamp=...	value=seppe.vandenbroucke@kuleuven.be
seppe	column=name:first,	timestamp=...	value=Seppe
seppe	column=name:firsttt,	timestamp=...	value=Seppe
seppe	column=name:last,	timestamp=...	value=vanden Broucke

HBase



● ● ● Fetching A HBASE Data By Row

```
get 'users', 'seppe'
```

● ● ● Fetching A HBASE Data By Row : Output

COLUMN

email:work

name:first

name:last

CELL

timestamp=1495293082872, value=seppe.vandenbroucke@kuleuven.be

timestamp=1495293050816, value=Seppe

timestamp=1495293067245, value=vanden Broucke

HBase



- HBase's query facilities are very limited
- Essentially a key-value, distributed data store with simple get/put operations
- Includes facilities to write MapReduce programs
- HBase (similar to Hadoop) doesn't perform well on less than 5 HDFS DataNodes with an additional NameNode
 - only makes the effort worthwhile when you can invest in, set up and maintain at least 6-10 nodes

Pig

- Yahoo! Developed “Pig”, which was made open source as Apache Pig in 2007
- High-level platform for creating programs that run on Hadoop which uses MapReduce underneath
 - The language used is **Pig Latin**
- Resembles the querying facilities of SQL



Pig



Loading a file using PIG LATIN

```
timesheet = LOAD 'timesheet.csv' USING PigStorage(',');
```



Filtering column 0

```
raw_timesheet = FILTER timesheet BY $0 > 100;
```



Fetching value based on alias

```
timesheet_logged = FOREACH raw_timesheet GENERATE $0 AS driverId, $2 AS hours_logged, $3 AS  
miles_logged;
```

Pig



Group BY

```
grp_logged = GROUP timesheet_logged BY driverId;
```



Group BY

```
sum_logged = FOREACH grp_logged GENERATE  
    group AS driverId,  
    SUM(timesheet_logged.hours_logged) AS sum_hourslogged,  
    SUM(timesheet_logged.miles_logged) AS sum_mileslogged;
```

Pig



- RDBMS and SQL are substantially faster than MapReduce – and hence Pig
- Pig Latin is relatively procedural versus declarative SQL

```
You tell the system how to do something, step-by-step.
```




Break

Hive



- Initially developed by Facebook but open-sourced afterwards
- Data warehouse solution offering SQL querying facilities on top of Hadoop
- Converts SQL-like queries to a MapReduce pipeline
- Also offers a JDBC and ODBC interface
- Can run on top of HDFS, as well as other file systems

Hive



- Hive **Metastore** stores metadata for each table such as its schema and location on HDFS.
- Driver service is responsible to receive and handle incoming queries
 - query is first converted to an abstract syntax tree, which is then converted to a directed acyclic graph representing an execution plan
 - the directed acyclic graph will contain a number of MapReduce stages and tasks

HQL Execution

Query



Driver Service



Abstract Syntax Tree (AST)



Directed Acyclic Graph (DAG)



MapReduce Stages

Hive



- HiveQL does not completely follow the full SQL-92 standard
 - E.g., lacks strong support for indexes, transactions, materialized views, and only has limited subquery support
- Example:
`SELECT genre, SUM(nrPages) FROM books
GROUP BY genre`
- HiveQL also allows to query data sets other than structured tables

Hive



```
CREATE TABLE docs (line STRING); -- create a docs table
```

```
-- load in file from HDFS to docs table, overwrite existing data:
```

```
LOAD DATA INPATH '/testfile.txt' OVERWRITE INTO TABLE docs;
```

```
-- perform word count
```

```
SELECT word, count(1) AS count
```

```
FROM ( -- split each line in docs into words
```

```
  SELECT explode(split(line, '\s')) AS word FROM docs
```

```
) GROUP BY t.word
```

```
ORDER BY t.word;
```

Hive



- One difference with traditional RDBMS is that Hive does not enforce the schema at the time of loading the data
 - Hive: schema-on-read
 - RDBMS: schema-on-write
- Recent versions of Hive support full ACID transaction management
- Performance and speed of SQL queries still forms the main disadvantage of Hive today
 - Solutions to bypass MapReduce (e.g. Apache Tez, Cloudera Impala, Facebook Presto)

Apache Spark




- Open-source alternative for MapReduce
- New programming paradigm centered on a data structure called the **Resilient Distributed Dataset (RDD)**
 - RDDs can enable the construction of iterative programs that have to visit a data set multiple times, as well as more interactive or exploratory programs
 - RDD is a fundamental data structure of Spark
 - Each dataset in RDD is divided into logical partitions that can be computed on different nodes of the cluster
 - Is maintained in a fault tolerant way
 - RDDs can contain any type of Python, Java or Scala objects, including user-defined classes.
- Many orders of magnitude faster than MapReduce implementations
- Rapidly adopted by many Big Data vendors

Apache Spark



- Similar to Hadoop, Spark works with HDFS and requires a cluster manager (e.g. YARN)
- Key components
 - Spark Core
 - Spark SQL
 - MLib, Spark Streaming, GraphX



Spark Core



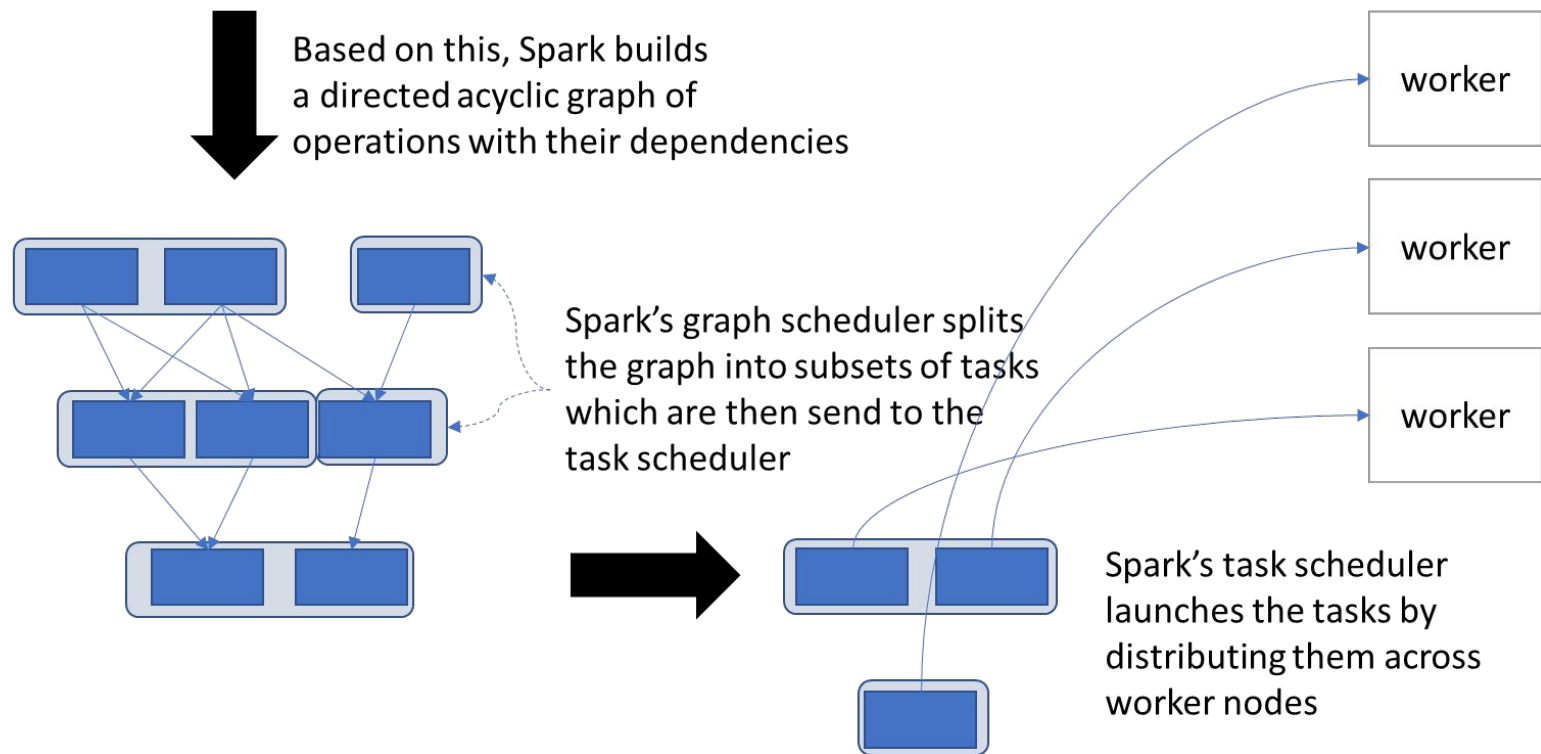
- Foundation for all other components
- Provides functionality for task scheduling and a set of basic data transformations
- Can be used through many programming languages
 - For example: Java, Python, Scala and R
- RDDs are the primary data abstraction in Spark
 - designed to support in-memory data storage and operations, distributed across a cluster
- Can be used to handle JSON and CSV data

Spark Core



A programmer writes a Spark program using its API:

```
rdd1.join(rdd2).groupBy(...).filter(...)
```



Spark Core



- Once data is loaded into an RDD, two basic types of operations can be performed:
 - Transformation which creates a new RDD through changing the original one
 - Actions which measure but do not change the original data
- Transformations are lazily evaluated
 - executed when a subsequent action has a need for the result
 - So can mean errors do not immediately appear
 - E.g., file does not exist
- RDDs will also be kept as long as possible in memory
- A chain of RDD operations gets compiled by Spark into a directed acyclic graph but which is then spread out and calculated over the cluster

Spark SQL



- Spark SQL runs on top of Spark Core and introduces another data abstraction called **DataFrames**
- DataFrames can be created from RDDs by specifying a schema on how to structure the data elements in the RDD, or can be loaded in directly from various sorts of file formats
- Even although DataFrames continue to use RDDs behind the scenes, they represent themselves to the end user as a collection of data organized into named columns
- You can run Spark directly using:
 - **pyspark** – uses Python
 - **spark-shell** – uses Scala
 - **spark-sql** – to run SQL queries
 - **spark-submit** – to run a program file, such as Python
- Or can access it via a programming language, such as Java
- See here for information:
 - <https://spark.apache.org/docs/latest/sql-programming-guide.html>

Spark Core – Map Reduce



- Spark's RDD API is relatively easy to work with compared to writing MapReduce programs

```
# Load in an RDD from a text file, the RDD will represent a  
# collection of text strings  
#(one for each line)
```

```
text_file = sc.textFile("testfile.txt")
```

Count the word occurrences

```
counts = text_file.flatMap(lambda line: line.split(" ")) \  
    .map(lambda word: (word, 1)) \  
    .reduceByKey(lambda a, b: a + b)
```

counts is a PythonRDD, Need loop to print items:

```
for x in counts.collect():  
    print x
```

```
(u'A', 1)  
(u'ago', 1)  
(u'episode', 1)  
(u'far', 2)  
(u'away', 1)  
(u'long', 1)  
(u'a', 1)  
(u'Another', 1)  
(u'Star', 1)  
(u'galaxy', 1)  
(u'of', 1)  
(u'in', 1)  
(u'Wars', 1)  
(u'time', 1)
```

Spark SQL – Handling JSON (pyspark)

Create a DataFrame object by reading in a file

```
df = spark.read.json("student.json")
```

```
df.show()
```

```
+---+-----+-----+---+
| age|      course|email| name|
+---+-----+-----+---+
| null|BSc Horticulture|    null|   Tom|
|  45| MSc Agriculture|    null| Helen|
|  30|      null|S.Carter.borchest...| Alice|
|  21|BSc Horticulture|    null|Johnny|
+---+-----+-----+---+
```

student.json:

```
{"name": "Tom",
  "course": "BSc Horticulture"}
{"name": "Helen", "age": 45,
  "course": "MSc Agriculture"}
{"name": "Alice", "age": 30,
  "email": "S.Carter@borchester.ac.uk"}
{"name": "Johnny", "age": 21,
  "course": "BSc Horticulture"}
```

DataFrames are structured in columns and rows:

```
df.printSchema()
```

```
root
 |-- age: long (nullable = true)
 |-- course: string (nullable = true)
 |-- email: string (nullable = true)
 |-- name: string (nullable = true)
```



Spark SQL (pyspark)



```
df.select("name").show()
```

```
+-----+  
|  name |  
+-----+  
|   Tom |  
|  Helen |  
|  Alice |  
|Johnny |  
+-----+
```

SQL-like operations can now easily be expressed:

```
df.select(df['name'], df['age'] + 1).show()
```

```
+-----+-----+  
|  name | (age + 1) |  
+-----+-----+  
|   Tom |      null |  
|  Helen |       46 |  
|  Alice |       31 |  
|Johnny |       22 |  
+-----+-----+
```

Spark SQL (pyspark)



```
df.filter(df['age'] > 21).show()
```

```
--+-----+-----+-----+
|age|      course|      email| name|
+---+-----+-----+-----+
| 45|MSc Agriculture|      null|Helen|
| 30|      null|S.Carter.borchest...|Alice|
+---+-----+-----+-----+
```

```
df.groupBy("course").count().show()
```

```
+-----+-----+
|      course|count|
+-----+-----+
| MSc Agriculture|    1|
|      null|    1|
|BSc Horticulture|    2|
+-----+-----+
```


Spark SQL (pyspark)



- Spark implements a full SQL query engine which can convert SQL statements to a series of RDD transformations and actions
- First Register the DataFrame as a SQL temporary view:
`df.createOrReplaceTempView("student")`

Can then use SQL like syntax:

```
sqlDF = spark.sql("SELECT name, age, course FROM student WHERE  
age > 21")
```

```
sqlDF.show()
```

```
+-----+-----+-----+  
| name|age|          course|  
+-----+-----+-----+  
|Helen| 45|MSc Agriculture|  
|Alice| 30|          null|  
+-----+-----+-----+
```

You can not just type
the SQL Code at the
command line

See Workbook for
further examples

MLlib, Spark Streaming and GraphX

- There are lots of other Spark tools to help work with Big Data:
 - **MLlib** is Spark's machine learning library
 - offers classification, regression, clustering, and recommender system algorithms
 - **Spark Streaming** uses Spark Core and its scheduling engine to perform streaming analytics
 - provides a high-level concept called **DStream**, which represents a continuous stream of data
 - **GraphX** is Spark's component implementing programming abstractions to deal with graph based structures
 - based on the RDD abstraction

MLlib, Spark Streaming and GraphX

■ Example: Word counting



```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext
sc = SparkContext("local[2]", "StreamingWordCount")
ssc = StreamingContext(sc, 1)

# Create a DStream that will connect to server.mycorp.com:9999 as a source
lines = ssc.socketTextStream("server.mycorp.com", 9999)

# Split each line into words
words = lines.flatMap(lambda line: line.split(" "))

# Count each word in each batch
pairs = words.map(lambda word: (word, 1))
wordCounts = pairs.reduceByKey(lambda x, y: x + y)

# Print out first ten elements of each RDD generated in the wordCounts Dstream
wordCounts.pprint()

# Start the computation
ssc.start()
ssc.awaitTermination()
```

Conclusion

- This lecture has looked at:
 - SQL on Hadoop
 - Apache Spark
- This week's workbook will look at using:
 - Further Hadoop examples
 - How to use CSV files
 - Apache Spark
 - SQL queries
 - JSON and CSV data